

CURSOR & MASTRA WORKSHOP

Mastra Studio walkthrough

What to do, and what you should see, at every checkpoint of the 90-minute build.

Start after setup is finished and your API keys are in place.

Studio walkthrough

Every checkpoint in the workshop, as a picture. Use this when you want to know what "working" is supposed to look like before you go hunting for a bug, or when you are facilitating and want the demo laid out in order.

This is a companion to stages 01 to 05, plus the stage 06 extension that puts the workflow in the agent's hands. It is not a replacement: the stages tell you what to build; this tells you what you should see afterwards.

Two things to know about the screenshots. The stage 01 images come from `before/`, so they show the real starting point. Everything from stage 02 onwards comes from the finished `after/` project, because that is the only place all the pieces exist at once — so its sidebar is fuller than yours will be at that moment. Judge the panel you are working in, not the item count next to it.

Getting it running

Assuming *setup* is done and your API keys are in place, three commands get you to the first screenshot.

1. Check the environment. From the repository root:

```
npm run verify
```

Nine `PASS` lines, ending in `All checks passed`. The one that matters most here is `MCP services respond 7 tools across 3 services`: it starts all three mock services, asks each for its tool list, and shuts them down again. If this fails, fix it before opening Studio — every stage below depends on those services answering.

You do not strictly have to run this. Each mock service creates and seeds its own database when it starts, and the policy service builds its vector index if it is missing, so a dev server can come up from nothing. Running `verify` first is still the better habit: it tells you *which* piece is broken, where a failure at boot just makes Studio look empty.

2. Start Studio on the project you are building in:

```
npm run dev          # the before/ project, where you work
npm run dev:after     # the finished reference, for comparison
```

3. Open the URL from the startup banner, not the one you remember:


```
| Studio: http://localhost:4111
```

Port 4111 is the default, but if something already has it, Mastra takes the next free port and prints that instead. Running `before/` and `after/` side by side is fine and normal — the second one just lands on 4112.

Leave the server running. Each stage below tells you when to restart it, which you need after adding anything new to `src/mastra/index.ts`, since a new top-level key does not always hot-reload.


Stage 01 — The empty starting point

Do this: `cd before && npm run dev` , then open the URL from the startup banner (4111 unless something else already has it).




Look for: one agent, `Setup Check` , and dashes in the Workflows, Agents, and Tools columns. That is not a broken install. Studio renders exactly what `src/mastra/index.ts` registers, and right now that file registers one throwaway agent.

Click **MCP Servers** in the sidebar and you get the same message in stronger terms:



The three services exist on disk in `src/mcp-servers/`, and `npm run verify` will confirm all seven of their tools answer. Studio still cannot see them, because nothing in your code has connected to them yet. Filling this sidebar is the whole workshop.


Prove your credentials work before building anything on top of them. Open **Agents** → **Setup Check** and say hello:



A reply means your API key, base URL, and model id are all good. The "Memory not enabled" notice on the left is expected: this agent has no memory, so nothing is saved as a thread.

Stage 02 — Your first tool


Do this: build the approval-route tool, register it, restart the dev server, and open **Tools**.



Look for: your tool listed under its `id`, `calculateApprovalRoute`, with the description you wrote next to it. At this point in the workshop it will be the only row; the `employee_`, `policy_`, and `expenseTool_` entries above arrive in stage 03.

That description column is worth a second of attention. It is not documentation for you — it is what the model reads when deciding whether to call your tool. A vague description is how a perfectly correct tool ends up never being used.

Then run it. Click the tool, enter `50`, and press Submit:




Look for: `approvers: ["manager"]` , `autoApproved: false` , `receiptRequired: true` .

Fifty is the interesting number. It is exactly on the boundary, so a tool written with `> 50` instead of `>= 50` auto-approves it and looks fine on every other input you try. Test `20` , `50` , `500` , and `640` : they should give you no approvers, manager, manager, and manager plus finance.


Stage 03 — The three services

Do this: build the MCP client, register the proxies, restart, and open **MCP Servers**.



Look for: three rows named `employee` , `policy` , and `expenseTool` — the keys you gave them in your client — with tool counts of 2, 1, and 4. Seven tools across three services is the same number `npm run verify` reports, so if the terminal and the browser disagree, trust neither and restart the dev server.

Then call one by hand. Click `expenseTool` → `listExpenses` and pass `emp-002` :




Look for: two historical claims, one `pending_approval` and one `approved`. You just ran code in a different process, over a protocol, with no agent involved anywhere. That separation is the entire point of MCP.

Notice the shape of the result: `{ result: { expenses: [...] } }`. MCP requires structured output to be an object, so the array is wrapped in one. A tool that returns a bare array fails with `expected record, received array`.

Stage 04 — The agent decides


Do this: build the citation skill and the agent, restart, and open **Agents** → **Expense Assistant**. Ask a policy question first:



Look for three things, in this order of importance:

1. **The citation.** [Meals & Entertainment Policy] after the claim. No citation means the skill is not loading — check that its description sounds relevant to the question.
2. The **Skill expense-policy-citations** chip above the answer. That is the agent deciding to load the skill. It is not always shown, and its absence is not a failure by itself, but the citation is.
3. **Policy searchExpensePolicy**. It looked the answer up rather than recalling it.

Then ask the question that shows what an agent actually is:



Look for: a chain of calls it chose on its own — the employee directory, your local `CalculateApprovalRoute`, and a policy search — and an answer that combines all three, with citations on each policy claim.

You never told it that order. You gave it capabilities and a goal, and it picked the sequence. Ask a neighbour to compare their trace with yours; they will differ, and that difference is the clearest definition of "agent" anyone will give you today.

Worth noticing in this answer: the \$500 finance threshold came from your **tool**, while the "equipment over \$200 needs IT approval" rule came from **policy search**. Deterministic arithmetic and prose knowledge, arriving through different doors, merged into one answer.

Stage 05 — The workflow pauses

Do this: build the workflow, restart, open **Workflows** → **expense-workflow**, and run it with `emp-002` and `Team dinner with 4 people in Chicago, $180 total`.


The screenshot displays the Mastra Studio interface for the 'expense-workflow'. The workflow is currently in a 'Suspended' state. The left sidebar shows the navigation menu with options like Primitives, Agents, Workflows, Processors, MCP Servers, Tools, Workspaces, Request Context, Evaluation, Overview, Scorers, Datasets, Experiments, and Observability. The main panel shows the workflow steps: 'request into a structured claim', 'route-for-approval', 'submit-claim', and 'human-approval'. The 'human-approval' step is highlighted with a 'Step suspended' panel on the right, which includes a 'Needs input' button and a 'Resume' button. A timeline at the bottom shows the duration of each step: Draft Claim (14.7s), Route For Approval (0s), Submit Claim (0.017s), and Human Approval (2.15s). The 'Recent runs' section shows a list of workflow runs with their IDs and timestamps.

Look for: the run marked **Suspended**, the first three steps green, and `human-approval` highlighted with a **Step suspended / Needs input** panel on the right. The timeline at the bottom shows where the time went: `draft-claim` takes seconds because it calls the agent, while `route-for-approval` and `submit-claim` are near-instant because they are ordinary code.

Check the drafted claim before you resume. The agent will usually have flagged the \$30/person cap on internal team meals, and some runs draft the amount as the \$150 that is actually reimbursable rather than the \$180 charged. Nobody told it to do that.

Now the part that matters. Stop the dev server and start it again, then reopen the run from **Recent runs**. It is still suspended, exactly where it was. The snapshot lives in storage, not in your browser or in a chat session, which is the reason workflows exist.

Then resume it. Tick the approve box, add a note, and press Resume:



Look for: status **Success** and all four steps green through to **End** . Confirm the decision was actually recorded by calling **MCP Servers** → **expenseTool** → **getExpenseStatus** with the **expenseId** from the result: the status should now be **approved** .


That last check is the one worth doing properly. The workflow did not write to **expenses.db** itself; it asked the service that owns that data to record the decision. A version that opens the database directly produces an identical screenshot and is still wrong.

Stage 06 — A receipt starts the whole thing

This is the stage 06 extension, not part of the main build. Not everyone will get here in the time, so treat it as the demo rather than a checkpoint you have to hit.

Up to now the agent and the workflow have lived in separate tabs: you ran the workflow yourself, from a form. The extension gives the agent the workflow as one of its tools, and that changes who starts the process.


Do this: in the Expense Assistant chat, use **Add attachment** → **Add a local file**, attach ``assets/sample-receipt.png``, and send *"here's my receipt for a team dinner. Please submit it for emp-002."*



Look for four separate things in that one reply:

- **The receipt was read.** Nobody typed `$180.00`, a vendor or a guest count. The chat model handles images, so the picture is the input.
- **The tools it chose**, unprompted: `getEmployee`, `searchExpensePolicy`, `calculateApprovalRoute` — and the `expense-policy-citations` skill firing underneath.
- **The policy arithmetic.** It works out that a \$30/person cap over 4 guests is \$120, so \$180 is \$60 over, and says a human approver will make the final call.
- **It submits anyway.** That last part is deliberate, and it is the interesting design decision in this whole workshop. An agent that refuses to submit an over-cap claim feels responsible, but it has quietly taken over the decision the `human-approval` step exists to make. The reference agent is told to submit it and record the breach, because the approver decides. Yours will probably stop and ask until you tell it otherwise.

Then it calls the workflow, and Studio draws the run inline in the conversation. Scroll down and you can watch the same four steps you ran from a form in stage 05:



Look for: the graph rendered inside the message, and `human-approval` marked as waiting with the run stopped there. The chat is now suspended too — the pause travelled all the way out to whoever was talking to the agent. Resume it from **Workflows** → **Recent runs**, exactly as in stage 05, and the claim completes.

That is the version worth showing someone sceptical about all of this: a photo of a receipt goes in, a routed and policy-checked claim comes out, and it still stops for a human before anything is approved.